

project_imbalanced_data

July 13, 2021

#imports

```
[57]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
import matplotlib.pyplot as plt
import seaborn as sns
import keras
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers import Adam
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
import cv2
import os
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.neighbors import KNeighborsClassifier
from sklearn import svm
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import mean_squared_error
```

Mounted at /content/drive

#Imbalanced Data

#generate data

```
[5]: labels = ['pedestrian', 'regular']
img_size = 128

def get_data(data_dir):
    data = []
    for label in labels:
        path = os.path.join(data_dir, label)
        class_num = labels.index(label)
        for img in os.listdir(path):
```

```

        try:
            img_arr = cv2.imread(os.path.join(path, img))[...,:-1]
            ↪#convert BGR to RGB format
            resized_arr = cv2.resize(img_arr, (img_size, img_size)) #
            ↪Reshaping images to preferred size
            data.append([resized_arr, class_num])
        except Exception as e:
            print(e)
    return np.array(data)

```

```
[6]: data = get_data('/content/drive/My Drive/machine learning/data')
```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:16:
VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences
(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths
or shapes) is deprecated. If you meant to do this, you must specify
'dtype=object' when creating the ndarray
    app.launch_new_instance()

#split to train and test

```

```
[7]: train,test = train_test_split(data, test_size=0.33, random_state=42)
print(train.shape)
print(test.shape)
```

```

(1605, 2)
(792, 2)

```

```
#show the balance in the train
```

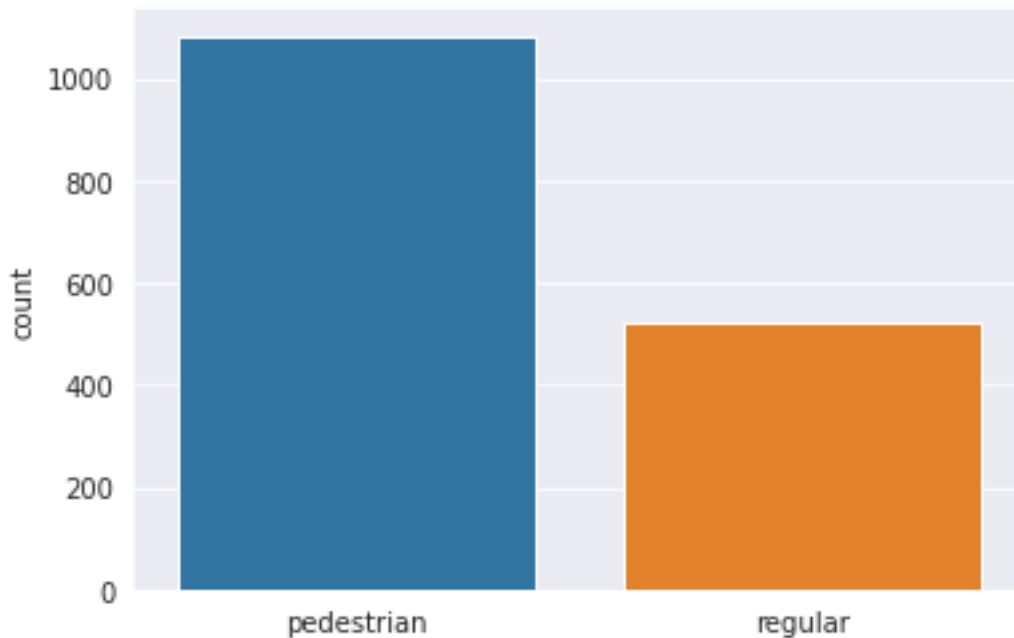
```
[8]: l = []
for i in train:
    if(i[1] == 0):
        l.append("pedestrian")
    else:
        l.append("regular")
sns.set_style('darkgrid')
sns.countplot(l)
```

```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
    FutureWarning

```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7f907e5923d0>
```



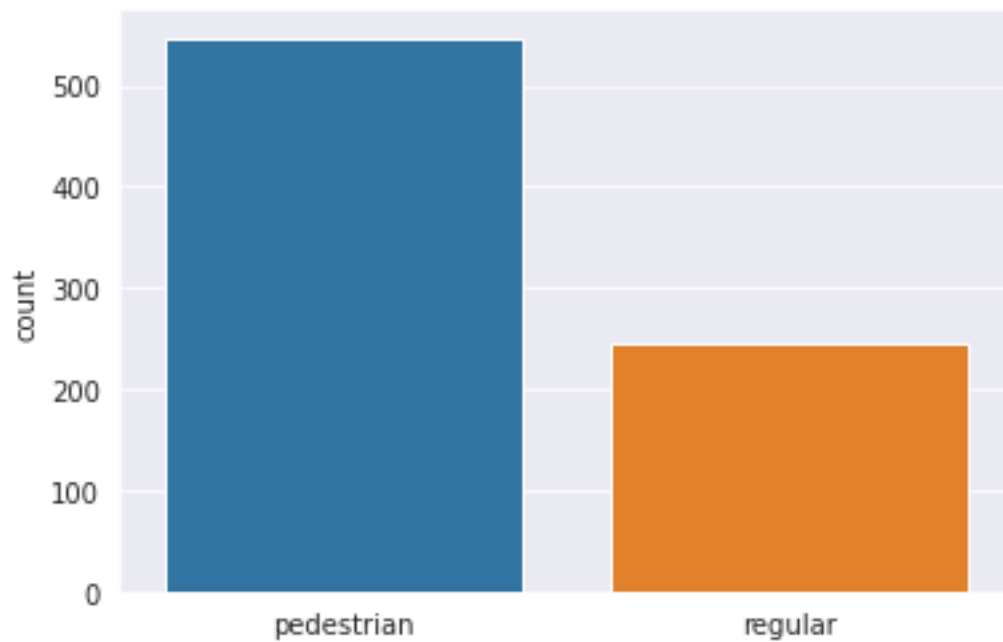
#show the balance in the test

```
[9]: r = []
    for i in test:
        if(i[1] == 0):
            r.append("pedestrian")
        else:
            r.append("regular")
    sns.set_style('darkgrid')
    sns.countplot(r)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning


```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f907e22a090>
```



#display some images from the data


```
[10]: plt.figure(figsize = (5,5))  
      plt.imshow(train[1][0])  
      plt.title(labels[train[0][1]])
```

```
[10]: Text(0.5, 1.0, 'pedestrian')
```



```
[11]: plt.figure(figsize = (5,5))  
      plt.imshow(train[-1][0])  
      plt.title(labels[train[-1][1]])
```

```
[11]: Text(0.5, 1.0, 'pedestrian')
```



#preper the data

```
[12]: x_train = []
      y_train = []
      x_test = []
      y_test = []

      for feature, label in train:
          x_train.append(feature)
          y_train.append(label)

      for feature, label in test:
          x_test.append(feature)
          y_test.append(label)

      # Normalize the data
      x_train = np.array(x_train) / 255
      x_test = np.array(x_test) / 255

      x_train.reshape(-1, img_size, img_size, 1)
      y_train = np.array(y_train)
```

```
x_test.reshape(-1, img_size, img_size, 1)
y_test = np.array(y_test)
```

```
[13]: print(x_train.shape)
      print(y_train.shape)
      print(x_test.shape)
      print(y_test.shape)
```

```
(1605, 128, 128, 3)
(1605,)
(792, 128, 128, 3)
(792,)
```

#do some random things on the data for better results

```
[14]: datagen = ImageDataGenerator(
        featurewise_center=False, # set input mean to 0 over the dataset
        samplewise_center=False, # set each sample mean to 0
        featurewise_std_normalization=False, # divide inputs by std of the
        ↪ dataset
        samplewise_std_normalization=False, # divide each input by its std
        zca_whitening=False, # apply ZCA whitening
        rotation_range = 30, # randomly rotate images in the range (degrees, 0
        ↪ to 180)
        zoom_range = 0.2, # Randomly zoom image
        width_shift_range=0.1, # randomly shift images horizontally (fraction
        ↪ of total width)
        height_shift_range=0.1, # randomly shift images vertically (fraction
        ↪ of total height)
        horizontal_flip = True, # randomly flip images
        vertical_flip=False) # randomly flip images

datagen.fit(x_train)
```

#CNN model

```
[15]: model = Sequential()
      model.add(Conv2D(32,3,padding="same", activation="relu",
        ↪ input_shape=(img_size,img_size,3)))
      model.add(MaxPool2D())

      model.add(Conv2D(32, 3, padding="same", activation="relu"))
      model.add(MaxPool2D())

      model.add(Conv2D(64, 3, padding="same", activation="relu"))
      model.add(MaxPool2D())
```

```

model.add(Dropout(0.6))

model.add(Flatten())
model.add(Dense(128,activation="relu"))
model.add(Dense(2, activation="softmax"))

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 32)	896
max_pooling2d (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_1 (Conv2D)	(None, 64, 64, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_2 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0
flatten (Flatten)	(None, 16384)	0
dense (Dense)	(None, 128)	2097280
dense_1 (Dense)	(None, 2)	258

Total params: 2,126,178
 Trainable params: 2,126,178
 Non-trainable params: 0

```

[16]: opt = Adam(lr=0.0001)
      model.compile(optimizer = opt , loss = tf.keras.losses.
      ↪SparseCategoricalCrossentropy(from_logits=True) , metrics = ['accuracy'])

```

```

/usr/local/lib/python3.7/dist-
packages/tensorflow/python/keras/optimizer_v2/optimizer_v2.py:375: UserWarning:
The `lr` argument is deprecated, use `learning_rate` instead.
  "The `lr` argument is deprecated, use `learning_rate` instead.")

```

```

[18]: epochs = 15

```



```
history = model.fit(x_train,y_train,epochs = epochs , validation_data =(x_test, y_test))
```

Epoch 1/15

51/51 [=====] - 40s 789ms/step - loss: 0.1134 - accuracy: 0.9620 - val_loss: 0.0815 - val_accuracy: 0.9760

Epoch 2/15

51/51 [=====] - 41s 796ms/step - loss: 0.0971 - accuracy: 0.9670 - val_loss: 0.0659 - val_accuracy: 0.9823

Epoch 3/15

51/51 [=====] - 40s 784ms/step - loss: 0.0704 - accuracy: 0.9769 - val_loss: 0.0482 - val_accuracy: 0.9886

Epoch 4/15

51/51 [=====] - 40s 784ms/step - loss: 0.0508 - accuracy: 0.9869 - val_loss: 0.0427 - val_accuracy: 0.9874

Epoch 5/15

51/51 [=====] - 41s 805ms/step - loss: 0.0438 - accuracy: 0.9850 - val_loss: 0.0331 - val_accuracy: 0.9937

Epoch 6/15

51/51 [=====] - 39s 775ms/step - loss: 0.0377 - accuracy: 0.9919 - val_loss: 0.0295 - val_accuracy: 0.9924

Epoch 7/15

51/51 [=====] - 41s 796ms/step - loss: 0.0290 - accuracy: 0.9938 - val_loss: 0.0232 - val_accuracy: 0.9962

Epoch 8/15

51/51 [=====] - 40s 785ms/step - loss: 0.0265 - accuracy: 0.9944 - val_loss: 0.0214 - val_accuracy: 0.9949

Epoch 9/15

51/51 [=====] - 40s 778ms/step - loss: 0.0223 - accuracy: 0.9944 - val_loss: 0.0168 - val_accuracy: 0.9962

Epoch 10/15

51/51 [=====] - 40s 786ms/step - loss: 0.0135 - accuracy: 0.9988 - val_loss: 0.0153 - val_accuracy: 0.9962

Epoch 11/15

51/51 [=====] - 40s 779ms/step - loss: 0.0152 - accuracy: 0.9963 - val_loss: 0.0192 - val_accuracy: 0.9937

Epoch 12/15

51/51 [=====] - 40s 778ms/step - loss: 0.0137 - accuracy: 0.9981 - val_loss: 0.0201 - val_accuracy: 0.9949

Epoch 13/15

51/51 [=====] - 39s 774ms/step - loss: 0.0143 - accuracy: 0.9969 - val_loss: 0.0148 - val_accuracy: 0.9975

Epoch 14/15

51/51 [=====] - 40s 777ms/step - loss: 0.0104 - accuracy: 0.9981 - val_loss: 0.0119 - val_accuracy: 0.9962

Epoch 15/15

51/51 [=====] - 40s 780ms/step - loss: 0.0095 - accuracy: 0.9981 - val_loss: 0.0134 - val_accuracy: 0.9962

#show results of CNN

```
[19]: acc = history.history['accuracy']
      val_acc = history.history['val_accuracy']
      loss = history.history['loss']
      val_loss = history.history['val_loss']

      epochs_range = range(epochs)

      plt.figure(figsize=(15, 15))
      plt.subplot(2, 2, 1)
      plt.plot(epochs_range, acc, label='Training Accuracy')
      plt.plot(epochs_range, val_acc, label='Validation Accuracy')
      plt.legend(loc='lower right')
      plt.title('Training and Validation Accuracy')

      plt.subplot(2, 2, 2)
      plt.plot(epochs_range, loss, label='Training Loss')
      plt.plot(epochs_range, val_loss, label='Validation Loss')
      plt.legend(loc='upper right')
      plt.title('Training and Validation Loss')
      plt.show()
```



```
[20]: predictions = model.predict_classes(x_test)
      predictions = predictions.reshape(1,-1)[0]
      print(classification_report(y_test, predictions, target_names = ['pedestrian_
      ↪(Class 0)','regular (Class 1)']))
```

/usr/local/lib/python3.7/dist-packages/keras/engine/sequential.py:450:
UserWarning: `model.predict_classes()` is deprecated and will be removed after

2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

warnings.warn("`model.predict_classes()` is deprecated and '

	precision	recall	f1-score	support
pedestrian (Class 0)	1.00	1.00	1.00	547
regular (Class 1)	0.99	1.00	0.99	245
accuracy			1.00	792
macro avg	1.00	1.00	1.00	792
weighted avg	1.00	1.00	1.00	792

#KNN model

```
[21]: x_train = x_train.reshape(x_train.shape[0],x_train.shape[1]*x_train.
      ↪shape[2]*x_train.shape[3])
      x_train.shape
```

[21]: (1605, 49152)

```
[22]: x_test = x_test.reshape(x_test.shape[0],x_test.shape[1]*x_test.shape[2]*x_test.
      ↪shape[3])
      x_test.shape
```

[22]: (792, 49152)

```
[23]: knn_1 = KNeighborsClassifier(n_neighbors=1)
      knn_1.fit(x_train,y_train)
      knn_2 = KNeighborsClassifier(n_neighbors=2)
      knn_2.fit(x_train,y_train)
      knn_3 = KNeighborsClassifier(n_neighbors=3)
      knn_3.fit(x_train,y_train)
      knn_4 = KNeighborsClassifier(n_neighbors=4)
      knn_4.fit(x_train,y_train)
      knn_5 = KNeighborsClassifier(n_neighbors=5)
      knn_5.fit(x_train,y_train)
      # knn_9 = KNeighborsClassifier(n_neighbors=9)
      # knn_9.fit(x_train,y_train)
```

```
[23]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
      metric_params=None, n_jobs=None, n_neighbors=5, p=2,
      weights='uniform')
```


#show results of KNN

```
[24]: score1 = knn_1.score(x_test, y_test)
score2 = knn_2.score(x_test, y_test)
score3 = knn_3.score(x_test, y_test)
score4 = knn_4.score(x_test, y_test)
score5 = knn_5.score(x_test, y_test)
print("1 neighbour: " ,score1)
print("2 neighbours: " ,score2)
print("3 neighbours: " ,score3)
print("4 neighbours: " ,score4)
print("5 neighbours: " ,score5)
```

```
1 neighbour: 0.9785353535353535
2 neighbours: 0.9406565656565656
3 neighbours: 0.946969696969697
4 neighbours: 0.9154040404040404
5 neighbours: 0.9356060606060606
```

```
[25]: plt.scatter(x_test[:,1],knn_1.predict_proba(x_test)[:,:1])
```

```
[25]: <matplotlib.collections.PathCollection at 0x7f90774c2e50>
```



#Logistic regression Model

```
[26]: logisticRegr = LogisticRegression()
result = logisticRegr.fit(x_train, y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940:
```

ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)


#show results of Logistic Regression

```
[27]: score = logisticRegr.score(x_test, y_test)
      print(score)
```

0.9873737373737373

```
[28]: plt.scatter(x_test[:,1],logisticRegr.predict_proba(x_test)[:,:1])
```

[28]: <matplotlib.collections.PathCollection at 0x7f90775cc610>



#SVM model

```
[29]: svm = svm.SVC(kernel='linear',probability=True) # Linear Kernel
      svm.fit(x_train, y_train)
```

```
[29]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
        decision_function_shape='ovr', degree=3, gamma='scale', kernel='linear',
        max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.001,
        verbose=False)
```


#show results of SVM

```
[30]: y_pred = svm.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
print("Precision:",metrics.precision_score(y_test, y_pred))
print("Recall:",metrics.recall_score(y_test, y_pred))
```

```
Accuracy: 0.9886363636363636
Precision: 0.9916666666666667
Recall: 0.9714285714285714
```

```
[31]: plt.scatter(x_test[:,1],svm.predict_proba(x_test)[:,:1])
```

```
[31]: <matplotlib.collections.PathCollection at 0x7f9077204210>
```



```
[32]: svm_preds = svm.predict(x_test)
knn_preds = knn_1.predict(x_test)
logistic_preds = logisticRegr.predict(x_test)
```

#Combine all models to see if it is improving results

```
[33]: # predictions
adaboost_preds = []

for x in range(len(predictions)):
    preds = []
    preds.append(svm_preds[x])
    preds.append(knn_preds[x])
    preds.append(logistic_preds[x])
    preds.append(predictions[x])
    adaboost_preds.append(max(set(preds), key = preds.count))
```

#results of the combined models

```
[34]: print("Accuracy:",metrics.accuracy_score(y_test, adaboost_preds))
print("Precision:",metrics.precision_score(y_test, adaboost_preds))
print("Recall:",metrics.recall_score(y_test, adaboost_preds))
```

Accuracy: 0.9924242424242424

Precision: 1.0

Recall: 0.9755102040816327

```
[35]: target_names = ['pedestrian', 'regular']
print(classification_report(y_test, adaboost_preds, target_names=target_names))
```

	precision	recall	f1-score	support
pedestrian	0.99	1.00	0.99	547
regular	1.00	0.98	0.99	245
accuracy			0.99	792
macro avg	0.99	0.99	0.99	792
weighted avg	0.99	0.99	0.99	792

```
[36]: print(mean_squared_error(y_test, adaboost_preds, multioutput='raw_values'))
```


[0.00757576]

```
[37]: print(confusion_matrix(y_test, adaboost_preds))
```

```
[[547  0]
 [ 6 239]]
```

```
[38]: plt.scatter(x_test[:,1],adaboost_preds)
```

```
[38]: <matplotlib.collections.PathCollection at 0x7f90770c7c90>
```



#Balanced Data

#Generate Data

```
[39]: labels = ['pedestrian', 'regular']
img_size = 128

def get_data(data_dir):
    data = []
    for label in labels:
        path = os.path.join(data_dir, label)
        class_num = labels.index(label)
        for img in os.listdir(path):
            try:
                img_arr = cv2.imread(os.path.join(path, img))[...,:-1]
                ↪ #convert BGR to RGB format
                resized_arr = cv2.resize(img_arr, (img_size, img_size)) #
                ↪ Reshaping images to preferred size
                data.append([resized_arr, class_num])
            except Exception as e:
                print(e)
    return np.array(data)
```

```
[40]: data = get_data('/content/drive/My Drive/machine learning/balanced data')
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:16:

VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences

(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

```
app.launch_new_instance()
```

#Split to train and test

```
[41]: train,test = train_test_split(data, test_size=0.33, random_state=42)
      print(train.shape)
      print(test.shape)
```

```
(1027, 2)
```

```
(507, 2)
```


#Show the balance in the train

```
[42]: l = []
      for i in train:
          if(i[1] == 0):
              l.append("pedestrian")
          else:
              l.append("regular")
      sns.set_style('darkgrid')
      sns.countplot(l)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7f907754b250>
```



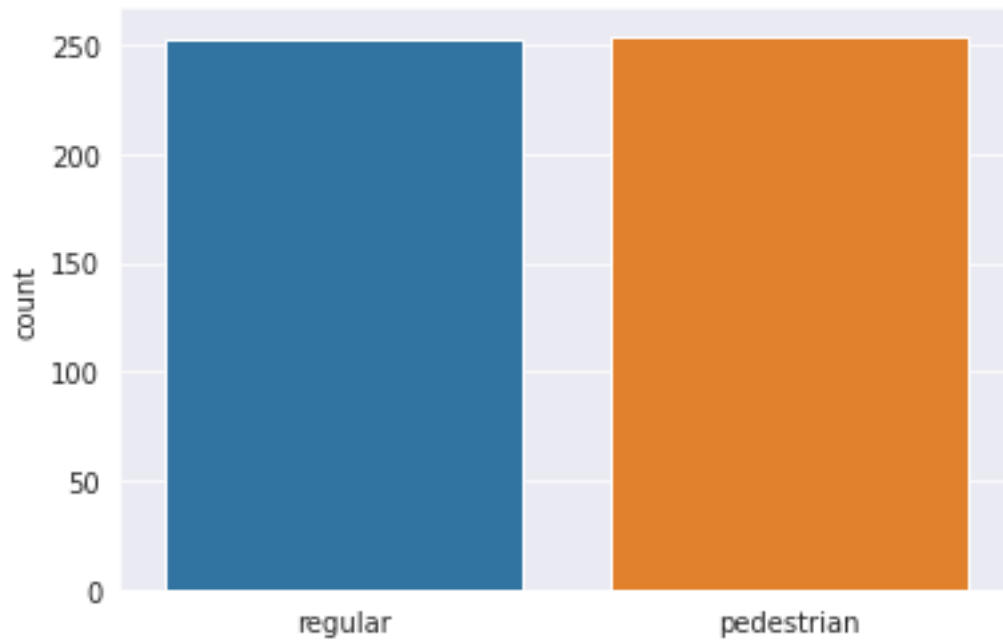
#Show the balance in the test

```
[43]: r = []
      for i in test:
          if(i[1] == 0):
              r.append("pedestrian")
          else:
              r.append("regular")
      sns.set_style('darkgrid')
      sns.countplot(r)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

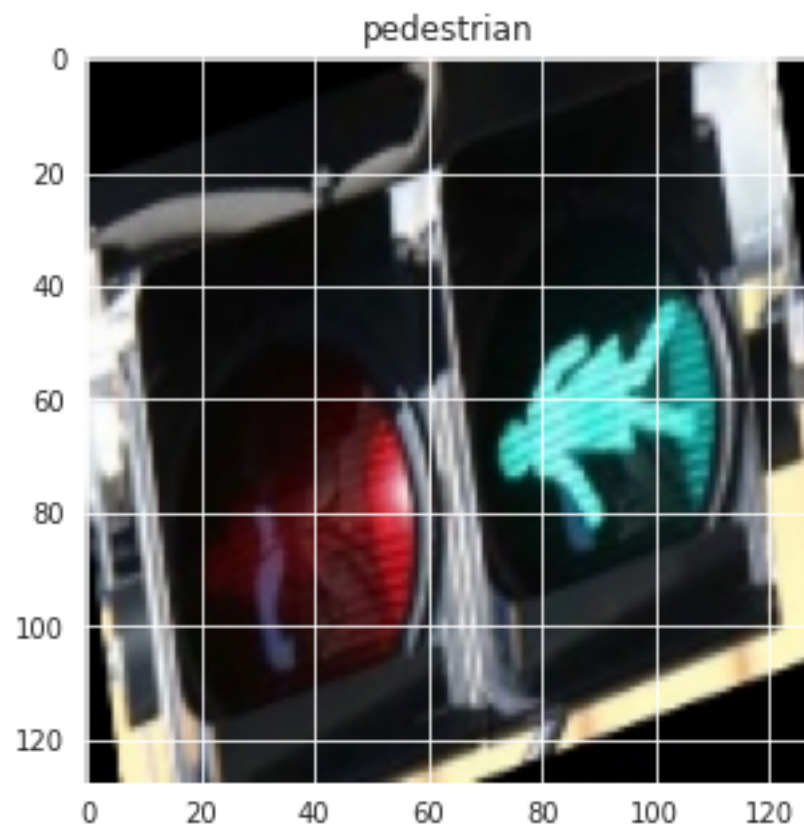
```
[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9073112210>
```

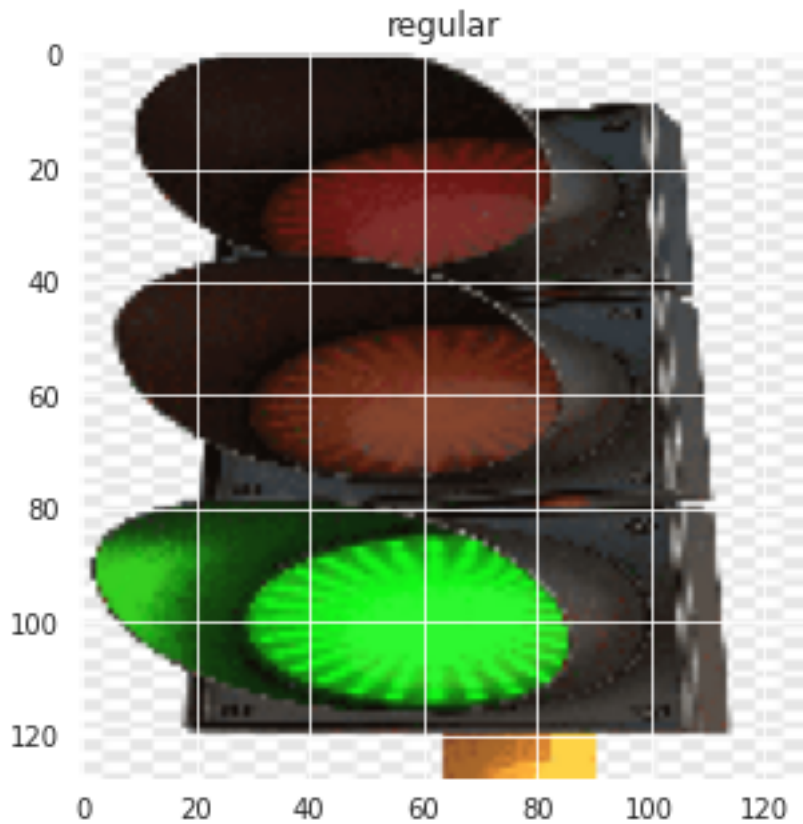


#Display some images from the data

```
[44]: plt.figure(figsize = (5,5))  
      plt.imshow(train[1][0])  
      plt.title(labels[train[0][1]])  
  
      plt.figure(figsize = (5,5))  
      plt.imshow(train[-1][0])  
      plt.title(labels[train[-1][1]])
```

```
[44]: Text(0.5, 1.0, 'regular')
```





#Preper the data

```
[45]: x_train_list = []
      y_train = []
      x_test_list = []
      y_test = []

      for feature, label in train:
          x_train_list.append(feature)
          y_train.append(label)

      for feature, label in test:
          x_test_list.append(feature)
          y_test.append(label)

      # Normalize the data
      x_train = np.array(x_train_list)/255
      x_test = np.array(x_test_list)/255
      # img = Image.fromarray(x_train[4])
      # img.save('my.png')
      # img.show()
```

```

x_train.reshape(-1, img_size, img_size, 1)
y_train = np.array(y_train)

x_test.reshape(-1, img_size, img_size, 1)
y_test = np.array(y_test)

print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

```

```

(1027, 128, 128, 3)
(1027,)
(507, 128, 128, 3)
(507,)

```

#Do some random things on the data for better results

```

[46]: datagen = ImageDataGenerator(
        featurewise_center=False, # set input mean to 0 over the dataset
        samplewise_center=False, # set each sample mean to 0
        featurewise_std_normalization=False, # divide inputs by std of the
        ↪ dataset
        samplewise_std_normalization=False, # divide each input by its std
        zca_whitening=False, # apply ZCA whitening
        rotation_range = 30, # randomly rotate images in the range (degrees, 0
        ↪ to 180)
        zoom_range = 0.2, # Randomly zoom image
        width_shift_range=0.1, # randomly shift images horizontally (fraction
        ↪ of total width)
        height_shift_range=0.1, # randomly shift images vertically (fraction
        ↪ of total height)
        horizontal_flip = True, # randomly flip images
        vertical_flip=False) # randomly flip images

datagen.fit(x_train)

```

#CNN model

```

[47]: class CNN:

        def __init__(self, x_train, y_train, x_test, y_test, img_size):
            self.x_train = x_train
            self.y_train = y_train
            self.x_test = x_test
            self.y_test = y_test
            self.img_size = img_size

```

```

self.model = self.model()

def model(self):
    model = Sequential()
    model.add(Conv2D(32, 3, padding="same", activation="relu",
↪input_shape=(self.img_size, self.img_size, 3)))
    model.add(MaxPool2D())

    model.add(Conv2D(32, 3, padding="same", activation="relu"))
    model.add(MaxPool2D())

    model.add(Conv2D(64, 3, padding="same", activation="relu"))
    model.add(MaxPool2D())
    model.add(Dropout(0.6))

    model.add(Flatten())
    model.add(Dense(128, activation="relu"))
    model.add(Dense(2, activation="softmax"))

    model.summary()
    return model

def run(self):

    opt = Adam(lr=0.0001)
    self.model.compile(optimizer=opt, loss=tf.keras.losses.
↪SparseCategoricalCrossentropy(from_logits=True),
                        metrics=['accuracy'])

    epochs = 15
    history = self.model.fit(self.x_train, self.y_train, epochs=epochs,
↪validation_data=(self.x_test, self.y_test))

    """#show results of CNN"""
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(epochs)

    plt.figure(figsize=(15, 15))
    plt.subplot(2, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy')
    plt.plot(epochs_range, val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
    plt.title('Training and Validation Accuracy')

```

```

plt.subplot(2, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

predictions = self.model.predict_classes(x_test)
predictions = predictions.reshape(1, -1)[0]
print(classification_report(self.y_test, predictions,
→target_names=['pedestrian (Class 0)', 'regular (Class 1)']))
return predictions

```

#KNN model

```

[48]: class KNN:

    def __init__(self, x_train, y_train, x_test, y_test):
        self.x_train = x_train
        self.y_train = y_train
        self.x_test = x_test
        self.y_test = y_test

    def run(self):
        knn_1 = KNeighborsClassifier(n_neighbors=1)
        knn_1.fit(self.x_train, self.y_train)
        knn_2 = KNeighborsClassifier(n_neighbors=2)
        knn_2.fit(self.x_train, self.y_train)
        knn_3 = KNeighborsClassifier(n_neighbors=3)
        knn_3.fit(self.x_train, self.y_train)
        knn_4 = KNeighborsClassifier(n_neighbors=4)
        knn_4.fit(self.x_train, self.y_train)
        knn_5 = KNeighborsClassifier(n_neighbors=5)
        knn_5.fit(self.x_train, self.y_train)

        """#show results of KNN"""

        score1 = knn_1.score(self.x_test, self.y_test)
        score2 = knn_2.score(self.x_test, self.y_test)
        score3 = knn_3.score(self.x_test, self.y_test)
        score4 = knn_4.score(self.x_test, self.y_test)
        score5 = knn_5.score(self.x_test, self.y_test)
        print("1 neighbour: ", score1)
        print("2 neighbours: ", score2)
        print("3 neighbours: ", score3)
        print("4 neighbours: ", score4)

```



```

print("5 neighbours: ", score5)

plt.scatter(self.x_test[:, 1], knn_1.predict_proba(self.x_test)[:, 1])
return knn_1

```

#Logistic Regression

```

[49]: class logisticRegression:

    def __init__(self, x_train, y_train, x_test, y_test):
        self.x_train = x_train
        self.y_train = y_train
        self.x_test = x_test
        self.y_test = y_test

    def run(self):
        logisticRegr = LogisticRegression()
        result = logisticRegr.fit(self.x_train, self.y_train)

        """#show results of Logistic Regression"""

        score = logisticRegr.score(self.x_test, self.y_test)
        print(score)

        plt.scatter(self.x_test[:, 1], logisticRegr.predict_proba(self.x_test)[:
→, 1])
        return logisticRegr

```

#SVM model

```

[50]: class SVM:

    def __init__(self, x_train, y_train, x_test, y_test):
        self.x_train = x_train
        self.y_train = y_train
        self.x_test = x_test
        self.y_test = y_test

    def run(self):
        SVM = svm.SVC(kernel='linear', probability=True) # Linear Kernel
        SVM.fit(self.x_train, self.y_train)

        """#show results of SVM"""

        y_pred = SVM.predict(self.x_test)
        print("Accuracy:", metrics.accuracy_score(self.y_test, y_pred))
        print("Precision:", metrics.precision_score(self.y_test, y_pred))
        print("Recall:", metrics.recall_score(self.y_test, y_pred))

```

```
plt.scatter(self.x_test[:, 1], SVM.predict_proba(self.x_test)[:, 1])
return SVM
```

#Adaboost model (combining all models)

```
[51]: class adaboost:

    def __init__(self, x_test, y_test, cnn, knn, lgr, svm):
        self.x_test = x_test
        self.y_test = y_test
        self.cnn = cnn
        self.knn = knn
        self.lgr = lgr
        self.svm = svm

    def run(self):
        cnn_preds = self.cnn
        cnn_acc = metrics.accuracy_score(self.y_test, cnn_preds)
        svm_preds = self.svm.predict(self.x_test)
        svm_acc = metrics.accuracy_score(self.y_test, svm_preds)
        knn_preds = self.knn.predict(self.x_test)
        knn_acc = metrics.accuracy_score(self.y_test, knn_preds)
        logistic_preds = self.lgr.predict(self.x_test)
        logistic_acc = metrics.accuracy_score(self.y_test, logistic_preds)

        # predictions
        adaboost_preds = []

        for x in range(len(self.cnn)):
            preds = 0.
            if svm_preds[x] == 0:
                preds += svm_acc * -1
            else:
                preds += svm_acc * 1
            if knn_preds[x] == 0:
                preds += knn_acc * -1
            else:
                preds += knn_acc * 1
            if logistic_preds[x] == 0:
                preds += logistic_acc * -1
            else:
                preds += logistic_acc * 1
            if cnn_preds[x] == 0:
                preds += cnn_acc * -1
            else:
                preds += cnn_acc * 1
```

```

        if preds >= 0:
            adaboost_preds.append(1)
        else:
            adaboost_preds.append(0)

        """#results of the combined models"""

        print("Accuracy:", metrics.accuracy_score(self.y_test, adaboost_preds))
        print("Precision:", metrics.precision_score(self.y_test,
→adaboost_preds))
        print("Recall:", metrics.recall_score(self.y_test, adaboost_preds))

        target_names = ['pedestrian', 'regular']
        print(classification_report(self.y_test, adaboost_preds,
→target_names=target_names))

        print(mean_squared_error(self.y_test, adaboost_preds,
→multioutput='raw_values'))

        print(confusion_matrix(self.y_test, adaboost_preds))

        plt.scatter(x_test[:, 1], adaboost_preds)

```

#Find Errors

```

[52]: def findErrors(x_test, y_test, algo, isCnn= False):
        lst = []
        if(isCnn):
            for idx, prediction, label in zip(enumerate(x_test), algo, y_test):
                if prediction != label:
                    lst.append(idx[0])
        else:
            for idx, prediction, label in zip(enumerate(x_test), algo.
→predict(x_test), y_test):
                if prediction != label:
                    lst.append(idx[0])
        # npArr = np.array(lst, dtype=int)
        print("list: ", lst)
        return np.array(lst, dtype=int)
        # np.save = ('sum.npy', npArr)

```

#Run models

#CNN Results

```

[53]: img_size = 128

```

```
print("\n CNN \n")
CNN = CNN(x_train, y_train, x_test, y_test, img_size).run()
```

CNN

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 128, 128, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_4 (Conv2D)	(None, 64, 64, 32)	9248
max_pooling2d_4 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_5 (Conv2D)	(None, 32, 32, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_1 (Dropout)	(None, 16, 16, 64)	0
flatten_1 (Flatten)	(None, 16384)	0
dense_2 (Dense)	(None, 128)	2097280
dense_3 (Dense)	(None, 2)	258

Total params: 2,126,178
 Trainable params: 2,126,178
 Non-trainable params: 0

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimizer_v2.py:375: UserWarning:
The `lr` argument is deprecated, use `learning_rate` instead.
"The `lr` argument is deprecated, use `learning_rate` instead.")
```

Epoch 1/15

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/backend.py:4930:
UserWarning: "`sparse_categorical_crossentropy` received `from_logits=True`, but
the `output` argument was produced by a sigmoid or softmax activation and thus
does not represent logits. Was this intended?"
```

```
"`sparse_categorical_crossentropy` received `from_logits=True`, but '
33/33 [=====] - 26s 762ms/step - loss: 0.6447 -
accuracy: 0.6136 - val_loss: 0.4559 - val_accuracy: 0.8047
```

Epoch 2/15
 33/33 [=====] - 25s 762ms/step - loss: 0.4132 - accuracy: 0.8057 - val_loss: 0.2872 - val_accuracy: 0.8876

Epoch 3/15
 33/33 [=====] - 25s 759ms/step - loss: 0.2323 - accuracy: 0.9364 - val_loss: 0.3566 - val_accuracy: 0.8146

Epoch 4/15
 33/33 [=====] - 25s 756ms/step - loss: 0.2655 - accuracy: 0.8803 - val_loss: 0.1607 - val_accuracy: 0.9566

Epoch 5/15
 33/33 [=====] - 25s 761ms/step - loss: 0.1205 - accuracy: 0.9659 - val_loss: 0.1321 - val_accuracy: 0.9586

Epoch 6/15
 33/33 [=====] - 25s 745ms/step - loss: 0.1221 - accuracy: 0.9529 - val_loss: 0.1168 - val_accuracy: 0.9665

Epoch 7/15
 33/33 [=====] - 24s 744ms/step - loss: 0.0890 - accuracy: 0.9772 - val_loss: 0.1029 - val_accuracy: 0.9704

Epoch 8/15
 33/33 [=====] - 25s 769ms/step - loss: 0.0855 - accuracy: 0.9718 - val_loss: 0.1028 - val_accuracy: 0.9724

Epoch 9/15
 33/33 [=====] - 26s 800ms/step - loss: 0.0746 - accuracy: 0.9796 - val_loss: 0.0845 - val_accuracy: 0.9744

Epoch 10/15
 33/33 [=====] - 28s 855ms/step - loss: 0.0652 - accuracy: 0.9867 - val_loss: 0.1081 - val_accuracy: 0.9448



Epoch 11/15
 33/33 [=====] - 27s 815ms/step - loss: 0.0619 - accuracy: 0.9747 - val_loss: 0.0722 - val_accuracy: 0.9822

Epoch 12/15
 33/33 [=====] - 26s 804ms/step - loss: 0.0647 - accuracy: 0.9713 - val_loss: 0.0874 - val_accuracy: 0.9763

Epoch 13/15
 33/33 [=====] - 27s 831ms/step - loss: 0.0549 - accuracy: 0.9808 - val_loss: 0.0756 - val_accuracy: 0.9684

Epoch 14/15
 33/33 [=====] - 26s 797ms/step - loss: 0.0339 - accuracy: 0.9923 - val_loss: 0.0753 - val_accuracy: 0.9645

Epoch 15/15
 33/33 [=====] - 26s 792ms/step - loss: 0.0243 - accuracy: 0.9957 - val_loss: 0.0697 - val_accuracy: 0.9783



```
/usr/local/lib/python3.7/dist-packages/keras/engine/sequential.py:450:
UserWarning: `model.predict_classes()` is deprecated and will be removed after
2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if
your model does multi-class classification (e.g. if it uses a `softmax` last-
layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model
does binary classification (e.g. if it uses a `sigmoid` last-layer
activation).
```

```
warnings.warn("`model.predict_classes()` is deprecated and '
```

	precision	recall	f1-score	support
pedestrian (Class 0)	0.97	0.99	0.98	254
regular (Class 1)	0.99	0.96	0.98	253
accuracy			0.98	507
macro avg	0.98	0.98	0.98	507
weighted avg	0.98	0.98	0.98	507

#KNN Results

```
[54]: x_train = x_train.reshape(x_train.shape[0], x_train.shape[1] * x_train.shape[2]
      ↪ * x_train.shape[3])
x_test = x_test.reshape(x_test.shape[0], x_test.shape[1] * x_test.shape[2] *
      ↪ x_test.shape[3])


print("\n KNN \n")
KNN = KNN(x_train, y_train, x_test, y_test).run()
```

KNN

```

1 neighbour: 0.9684418145956607
2 neighbours: 0.9250493096646942
3 neighbours: 0.9349112426035503
4 neighbours: 0.903353057199211
5 neighbours: 0.9230769230769231

```



#Logistic Regression Results

```

[55]: print("\n Logistic Regression \n")
      LGR = logisticRegression(x_train, y_train, x_test, y_test).run()

```

Logistic Regression

```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:


<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

0.9664694280078896




#SVM Results

```
[58]: print("\n SVM \n")
      SVM = SVM(x_train, y_train, x_test, y_test).run()
```

SVM

Accuracy: 0.9644970414201184
Precision: 0.9757085020242915
Recall: 0.9525691699604744



#Adaboost Results

```
[59]: print("\n Adaboost \n")
adb = adaboost(x_test, y_test, CNN, KNN, LGR, SVM)
adb.run()
```

Adaboost

Accuracy: 0.980276134122288

Precision: 0.9879518072289156


Recall: 0.9723320158102767

	precision	recall	f1-score	support
pedestrian	0.97	0.99	0.98	254
regular	0.99	0.97	0.98	253
accuracy			0.98	507
macro avg	0.98	0.98	0.98	507
weighted avg	0.98	0.98	0.98	507

[0.01972387]

[[251 3]

[7 246]]



```
[75]: new_data = get_data('/content/drive/My Drive/machine learning/new_data')
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:16:
VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences
(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths
or shapes) is deprecated. If you meant to do this, you must specify
'dtype=object' when creating the ndarray
app.launch_new_instance()
```

```
[85]: x = []
y = []

for feature, label in new_data:
    x.append(feature)
    y.append(label)

# Normalize the data
x = np.array(x) / 255

x.reshape(-1, img_size, img_size, 1)
y = np.array(y)
```

```
[87]: cnn_preds = model.predict_classes(x)
cnn_acc = metrics.accuracy_score(y, cnn_preds)
print(cnn_preds)
print(cnn_acc)
```

```
[1 1 0 0 0 1 0 1 1 0]
0.6
```

```
/usr/local/lib/python3.7/dist-packages/keras/engine/sequential.py:450:
UserWarning: `model.predict_classes()` is deprecated and will be removed after
2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if
your model does multi-class classification (e.g. if it uses a `softmax` last-
layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model
does binary classification (e.g. if it uses a `sigmoid` last-layer
activation).
```

```
warnings.warn("`model.predict_classes()` is deprecated and '
```

```
[94]: x_ = x.reshape(x.shape[0],x.shape[1]*x.shape[2]*x.shape[3])
      x_.shape
```

```
[94]: (10, 49152)
```

```
[95]: svm_preds = SVM.predict(x_)
      svm_acc = metrics.accuracy_score(y, svm_preds)
      knn_preds = KNN.predict(x_)
      knn_acc = metrics.accuracy_score(y, knn_preds)
      logistic_preds = LGR.predict(x_)
      logistic_acc = metrics.accuracy_score(y, logistic_preds)
```

```
[96]: print("SVM" , svm_acc)
      print("KNN" ,knn_acc)
      print("Logistic Regression" ,logistic_acc)
```

```
SVM 0.8
```

```
KNN 0.6
```

```
Logistic Regression 0.6
```

```
[97]: adb = adaboost(x_, y, cnn_preds, KNN, LGR, SVM)
      adb.run()
```

```
Accuracy: 0.7
```

```
Precision: 0.6666666666666666
```

```
Recall: 0.8
```

	precision	recall	f1-score	support
pedestrian	0.75	0.60	0.67	5
regular	0.67	0.80	0.73	5
accuracy			0.70	10
macro avg	0.71	0.70	0.70	10
weighted avg	0.71	0.70	0.70	10

```
[0.3]
```

```
[[3 2]
```

```
[1 4]]
```

```
↳ -----  
ValueError                                Traceback (most recent call↳  
↳last)  
  
    <ipython-input-97-5742f352dcf4> in <module>()  
        1 adb = adaboost(x_, y, cnn_preds, KNN, LGR, SVM)  
----> 2 adb.run()  
  
    <ipython-input-51-1c29ab7cadb8> in run(self)  
    59         print(confusion_matrix(self.y_test, adaboost_preds))  
    60  
--> 61         plt.scatter(x_test[:, 1], adaboost_preds)  
  
    /usr/local/lib/python3.7/dist-packages/matplotlib/pyplot.py in ↳  
↳scatter(x, y, s, c, marker, cmap, norm, vmin, vmax, alpha, linewidths, verts, ↳  
↳edgecolors, plotnonfinite, data, **kwargs)  
    2814         verts=verts, edgecolors=edgecolors,  
    2815         plotnonfinite=plotnonfinite, **({"data": data} if data is not  
-> 2816         None else {}), **kwargs)  
    2817     sci(__ret)  
    2818     return __ret  
  
    /usr/local/lib/python3.7/dist-packages/matplotlib/__init__.py in ↳  
↳inner(ax, data, *args, **kwargs)  
    1563     def inner(ax, *args, data=None, **kwargs):  
    1564         if data is None:  
-> 1565             return func(ax, *map(sanitize_sequence, args), **kwargs)  
    1566  
    1567         bound = new_sig.bind(ax, *args, **kwargs)  
  
    /usr/local/lib/python3.7/dist-packages/matplotlib/cbook/deprecation.py ↳  
↳in wrapper(*args, **kwargs)  
    356         f"%(removal)s.  If any parameter follows {name!r}, ↳  
↳they "  
    357         f"should be pass as keyword, not positionally.")  
--> 358         return func(*args, **kwargs)  
    359  
    360     return wrapper
```

```

/usr/local/lib/python3.7/dist-packages/matplotlib/axes/_axes.py in
↳ scatter(self, x, y, s, c, marker, cmap, norm, vmin, vmax, alpha, linewidths,
↳ verts, edgecolors, plotnonfinite, **kwargs)
    4389         y = np.ma.ravel(y)
    4390         if x.size != y.size:
-> 4391             raise ValueError("x and y must be the same size")
    4392
    4393         if s is None:

```

ValueError: x and y must be the same size

